

Guideline for AVeriTeC annotators

v1.0 - May 17, 2022

1 Introduction

We aim to construct a dataset for automated fact-checking with the following guiding principles. First, we intend to decompose the evidence retrieval process into multiple steps, annotating each individual step as a question-answer pair (see Figure 1). Second, our dataset will be constructed from real-world claims previously checked by journalistic organisations, rather than the artificially created claims used in prior work.

Decomposing claim verification into generation and answering of questions allows us to break complex real-world claims down to their components, simplifying the task. For example, in Figure 1, verifying the claim requires knowing the salary of the health commissioner, the governor, the vice president and Dr. Fauci, so that they can be compared. Four separate questions about salary need to be asked in order to reach a verdict (i.e. that the claim is *supported*).

By decomposing the evidence retrieval process in this way, we also produce a natural way for systems to justify their verdicts and explain their reasoning to users. In addition to this, we annotate claims with a final justification, providing a textual explanation of how to combine the retrieved answers to reach a verdict. This allows users to follow each step of the retrieval and verification processes, and so understand the reasoning employed by the system.

Claim: Biden lead disappears in NV, AZ, GA, PA on 11 November 2020.

Q1: Which media project Biden will win in Nevada?
A1: ABC News, CBS News, NBC News, CNN, Fox News, Decision Desk HQ, Associated Press, Reuters, and New York Times.

Q2: Which media project Biden will win in Arizona?
A2: Fox News and Associated Pre.

Q3: Which media project Biden will win in Georgia?
A3: None.

Q4: Which media project Biden will win in Pennsylvania?
A4: ABC News, CBS News, NBC News, CNN, Fox News, Decision Desk HQ, Associated Press, Reuters, and New York Times.

Verdict: Refuted
Justification: Many media organizations believe Biden will win in NV, AZ, and PA. As such, his lead has not disappeared.

Figure 1: Example claim and question answer pairs.

The annotation consists of the following three phases:

1. Claim Normalization.
2. Question Generation.
3. Quality Control.

Each claim should be annotated by different annotators in each phase. An annotator can participate *in* all three phases, but they will be assigned different claims.

Warning! Components of the AVeriTeC annotation tool may not render correctly in some browsers, specifically *Opera Mini*. If this is an issue we recommend trying another browser, e.g. Firefox, Chrome, Safari, or regular Opera.

2 Sign In

Each annotator will have received an **ID** and a **Password** with the access link to the annotation server. The password can be changed after logging into the interface.

Important!

- Make sure to log out at the end of the session!
- Do not open multiple tabs/windows of the AVeriTeC annotation tool. Always use only one window during annotation! If you are logged into multiple sessions using the same account, the annotation tool may lose data you enter.



Figure 2: Interface of the control panel. (1) Button for changing the password. (2) Button for logout. (3) Start the annotation for this phase. Here is Phase 1 Claim Normalization. (4) The left number shows how many claims have been annotated and the right number shows how many claims are assigned for the current annotator at this phase.

After clicking the **START NEXT** button, the annotation phase will start. If an annotator is new to the current phase, the interface will provide a guided tour as in Figure 3 for that phase. Please read the hints provided by the tour guide carefully before the annotation.

3 Phase 1: Claim Normalization

In the first phase, annotators collect metadata about the claims, and produce a normalized version of each claim, as shown in Figure 4. The first step is to identify the claim(s) in the fact-checking article. Often, this can be found either in the headline or explicitly some other place in the fact-checking article. In some cases, there may be a discrepancy between the article and the original claim (e.g. the original claim could be “*there are 30 days in March*”, while the fact-checking article might have the headline “*actually, there are 31 days in March*”). In those cases, it is important to use the *original* version of the claim. If there is ambiguity in the article over the exact wording of the claim, annotators should use their own judgment.

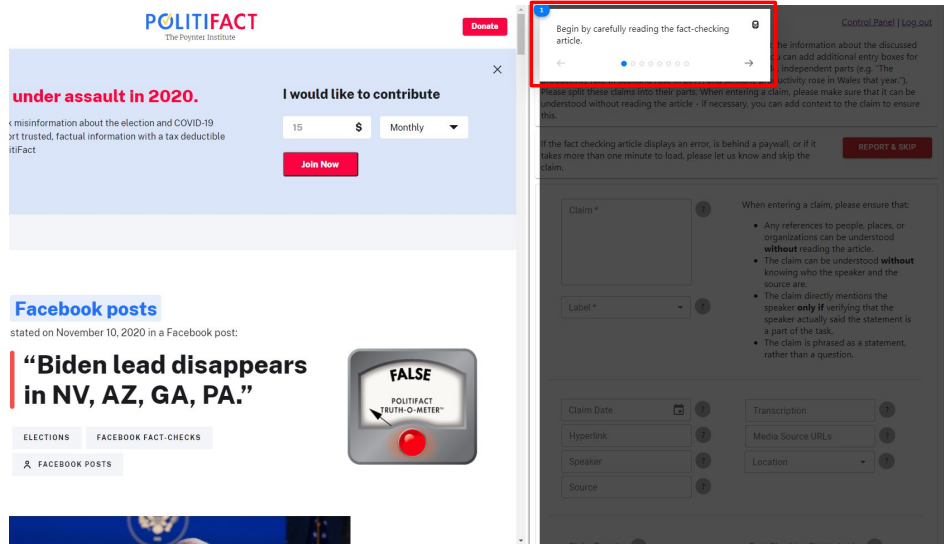


Figure 3: Interface of the tour guide.

Figure 4: Interface of claim normalization. ① The fact-checking article provided. ② Guideline of annotation for this phase. Please read it before annotating. Notice that if the article displays a 404 page or another error, or if it takes more than one minute to load, please click the **REPORT & SKIP** button. ③ Fields for the normalized claim and the corresponding label. ④ General information of the claim. ⑤ Check-boxes for selecting the type of the claim. ⑥ Check-boxes for selecting the fact-checking strategy used. ⑦ Button for adding more claims. ⑧ Buttons for submitting the current claim, going to the previous claim, and next claim.

3.1 Overview

Here, we give a quick overview of the claim normalization task; in-detail discussion can be found in subsequent sections. Further documentation can also be found on-the-fly using the tooltips in the annotation interface.

1. First, annotators should read the fact-checking article and identify which claims are being investigated.
2. If the fact-checking article is paywalled or inaccessible due to a 404-page or a similar error message, annotators should report this and skip the claim using the provided button. We warn that some fact-checking articles can take long too load – as such, while fact-checking articles that do not load at all should be skipped, we ask annotators to wait for at least one minute before skipping an article while it is still trying to load.
3. Most articles focus on one claim. However, some articles investigate multiple claims, or claims with multiple parts – in those cases, annotators should first split these into their parts (see Section 3.2).
4. Some claims cannot be understood without the context of the fact-checking article, e.g. because they refer to entities not mentioned by name in the claim. In those cases, annotators should add context to the claims (see Section 3.3).
5. Generally, we prefer claims to be as close as possible to their original form (i.e. the form originally said, *not* the form used in the fact-checking article). As such, contextualization should be done only when necessary, following the checklist in Section 3.3.
6. Annotators should extract the verdict assigned to the claim in the article, and translate it as closely as possible to one of our four labels – *supported*, *refuted*, *not enough evidence*, or *conflicting evidence/cherry picking* (see Section 3.4). In phase one, annotators should give their own judgments – rather, they should match as closely as possible the judgments given by the fact-checking articles.
7. Claims will have associated metadata, i.e. the date the original claim was made, or the name of the person who made it. Annotators should identify and extract this metadata from the article (see Section 3.5).
8. Annotators should identify the type of each claim, choosing from the options described in Section 3.5.2. These are not mutually exclusive, and more than one claim type can be chosen.
9. Annotators should identify the strategies used in the fact-checking article to verify each claim, choosing from the options described in Section 3.5.3. These are not mutually exclusive, and more than one claim type can be chosen.

3.2 Claim Splitting

Some claims consist of multiple, easily separable, independent parts (e.g. “*The productivity rate in Scotland rose in 2017, and similarly productivity rose in Wales that year.*”). The first step is to split these compound claims into individual claims. Metadata collection and normalization will then be done independently for each individual claim, and in subsequent phases they will be treated as separate claims.

When splitting a claim, it is important to ensure that each part is understandable without requiring the others as context. This can be done either by adding metadata in the appropriate field, such as the claim speaker or claim date, or through rewriting. For example, for the claim “*Amazon is doing great damage to tax paying retailers. Towns, cities and states throughout the U.S. are being hurt - many jobs being lost!*”, it should be clear what is causing job loss in the second part. A possible split would be “*Amazon is doing great damage to tax paying retailers*” and “*Towns, cities and states throughout the U.S. are being hurt by Amazon - many jobs being lost*”. That is, it is necessary to rewrite the second part by adding *Amazon* a second time in order for the second part to be understandable without context.

3.3 Claim Contextualization

Some claims are not complete, which means they lack adequate contextualization to be verified. For example, in the claim *“We have 21 million unemployed young men and women.”*, there are unresolved pronouns without which the claim cannot be verified (e.g. *we* refers to Nigeria, as the speaker of the claim is the presidential candidate of Nigeria). Another example is *“Israel already had 50% of its population vaccinated.”* We need to know when this claim was made to verify its veracity, as the time is crucial for this verification. For the latter, metadata is enough to resolve ambiguities; the former needs to be rewritten as *“Nigeria has 21 million unemployed young men and women.”*

Annotators are asked to contextualize claims to the original post by gathering the necessary information. Some information can be included simply as metadata, but this is not always enough – for information not captured by metadata, we ask that the claim itself is rewritten to include said information. Annotators need to follow this checklist:

1. Is the claim referring to entities which can only be identified by reading the associated fact-checking article, even if all metadata is taken into consideration? If so, add the names of the entities (e.g. *“Former first lady said, ‘White folks are what’s wrong with America’.”* becomes *“Former first lady Michelle Obama said, ‘White folks are what’s wrong with America’.”*).
2. Does the claim have unnecessary quotation marks, or references to a speaker (such as the word *says* in the example here)? If so, remove them (e.g. *“Says ‘Monica Lewinsky Found Dead’ in a burglary.”* becomes *“Monica Lewinsky found dead in a burglary.”*). Do NOT remove the reference to the speaker if the central problem is to determine if that person actually said the quote, e.g. in the case of quote verification.
3. Is the claim a question? If so, rephrase it as a statement (e.g. *“Did a Teamsters strike hinder aid efforts in Puerto Rico after Hurricane Maria?”* becomes *“A Teamsters strike hindered aid efforts in Puerto Rico after Hurricane Maria in 2017.”*).
4. Does the claim contain pronominal references to entities only mentioned in the fact-checking article? If so, replace the pronoun with the name of that entity. (e.g. *“We have 21 million unemployed young men and women.”* becomes *“Nigeria has 21 million unemployed young men and women.”*).
5. For some fact-checking articles, the title used does not properly match the fact-checked claim. Find the original claim in the article, and use that for producing the normalized version. As shown in Figure 5, the claim should be the first sentence of the article rather than the title.
6. Is the claim too vague to be investigated through the use of evidence, and does the fact-checking article investigate a more specific version of the claim? If so, use the claim investigated in the fact-checking article (e.g. *“Towns, cities and states throughout the U.S. are being hurt by Amazon”* might become *“Towns, cities and states throughout the U.S. are losing state tax revenue because of Amazon”*).

Generally, try to make claims specific enough so that they can *be understood* and so that *appropriate evidence can be found* by a person who has not seen the fact-checking article.

Important! We recommend reading through the entire article and understanding the central problem before rewriting the claim. This makes it easier to identify the exact phrasing of the original claim, and to make any minimal interventions necessary following our checklist above. When in doubt as to whether a claim should be modified, we recommend leaving it unchanged – we generally prefer claims to as close as possible to their original form, subject to the constraints listed above.

Hoax Alert

Fake News: Donald Trump Did NOT Say "Women, You Have To Treat Them Like Sh*t"

Title

Sep 16, 2019 by: Maarten Schenk

Share

Tweet

Did Donald Trump say "Women, you have to treat them like shit" in a 1992 interview in New York Magazine?

that's not true: Trump was not speaking about women general but about a very narrow category of people, to be particular "supermodels... clinging to a rock star's legs". While Trump has definitely insulted or disparaged many individual women in his career (and men... and groups of people...), he did not directly say *all* women have to be treated like shit in this interview.

The quote recently reappeared in [an image post](#) (archived [here](#)) published by Occupy Democrats on Facebook on September 14, 2019:

Claim

Figure 5: An example of locating the claim.

3.4 Labels

We ask annotators to produce a label for the claim relying *only* on the information on the fact-checking site (and assuming that everything reported there is accurate). For the dataset we are creating, we will be using four labels:

1. The claim is **supported**. The claim is supported by the arguments and evidence presented.
2. The claim is **refuted**. The claim is contradicted by the arguments and evidence presented.
3. There is **not enough evidence** to support or refute the claim. The evidence either directly argues that appropriate evidence cannot be found, or leaves some aspect of the claim neither supported nor refuted. We note that many fact-checking agencies mark claims as *refuted* (or similar), if supporting evidence does not exist, without giving any refuting evidence. We ask annotators to use *not enough evidence* for this category, regardless of the original label.
4. The claim is misleading due to **conflicting evidence/cherry-picking**, but not explicitly refuted. This includes cherry-picking, true-but-misleading claims (e.g. the claim "*Alice has never lost an election*" with evidence showing Alice has only ever run unopposed), as well as cases where conflicting or internally contradictory evidence can be found. Missing context may also be relevant if a situation has changed over time, and the claim fails to mention this (e.g. "*Alice is a strong supporter of industrial subsidies*" with evidence showing that Alice currently supports the position, but in the past opposed the position).

Despite the claim splitting subtask, some claims may contain multiple parts that are too interconnected to split. This could for example be a claim like "*Alice has never lost an election because she always supports cheese subsidies*". In such cases, parts of the claim may have different truth values. We discuss a few cases below:

- The claim is implicature, i.e. “*X happens because Y*” or “*X leads to Y*”. In this case, annotators should find a label for the causal implication, and *not* for either of the component claims.
- The claim has too components, where one is *refuted* and the other is *not enough information*. In this case, the entire claim should be labelled *refuted*.
- The claim has too components, where one is *supported* and the other is *not enough information*. In this case, the entire claim should be labelled *not enough information*.

Important! The label given in Phase 1 – and *only* in Phase 1 – should reflect the decision of the fact checker, not the interpretation of the annotator. In Phase 1, annotators should report the original judgment, as closely as possible, even if they disagree with it.

3.5 Metadata Collection

Annotators need to collect metadata through the following three steps.




Figure 6: An example of an image claim requiring transcription.

3.5.1 General Information

- A hyperlink to the original claim, if that is provided by the fact-checking site. Examples of this include Facebook posts, the original article or blog post being fact checked, and embedded video links. If the original claim has a hyperlink on the fact-checking site, but that hyperlink is dead, annotators should leave the field empty.
- The date of the original claim, regardless of whether it is necessary for verifying the claim. This date is often mentioned by the fact checker, but not in a standardized place where we could automatically retrieve it. Note that the date for the *original claim* and the *fact-checking article* (often its publication date) may be different and both stated in text. We specifically need the original claim date, as we intend to filter out evidence that appeared after that date. If multiple dates are mentioned, the earliest should be used. If an imprecise date is given (e.g. February 2017), the earliest possible interpretation should be used (i.e. February 1st, 2017).
- The speaker of the original claim, e.g. the person or organization who made the claim.

- The source of the original claim, e.g. the person or organization who published the claim. This is not necessarily the same as the speaker; a person might make a comment in a newspaper, in which case the person is the speaker and the newspaper is the source.
- If the original claim is or refers to an image, video, or audio file, annotators should add a link to that media file (or the page that contains the file, if the media file itself is inaccessible).
- If the original claim is an image that contains text – for example, Figure 6 shows a Facebook meme about Michelle Obama – annotators should transcribe the text that occurs in the image as metadata. In the example, it would be “*Michelle Obama said white folks are what’s wrong with America.*”
- If the fact-checking article is paywalled or inaccessible due to an error message, annotators should report this and skip the claim using the corresponding button.

3.5.2 Claim Type

The type of the claim itself, independent of the approach taken by the fact checker to verify or refute it, should be chosen from the following list. This is not a mutually exclusive choice – a claim can be speculation about a numerical fact, for example. As such, annotators should choose one *or several* from the list.

- **Speculative Claim:** The primary task is to assess whether a prediction is plausible or realistic. For example “*the price of crude oil will rise next year.*” The primary task is to assess whether the prediction is plausible or realistic.
- **Opinion Claim:** The claim is a non-factual opinion, e.g. “*cannabis should be legalized*”. This contrasts with factual claims on the same topic, such as “*legalization of cannabis has helped reduce opioid deaths.*”
- **Causal Claim:** The primary task is to assess whether one thing caused another. For example “*the price of crude oil rose because of the Suez blockage.*”
- **Numerical claim.** The primary task is to verify whether a numerical fact is true, or to verify whether a comparison between several numerical facts hold, or to determine whether a numerical trend or correlation is supported by evidence.
- **Quote Verification.** The primary task is to identify whether a quote was actually said by the supposed speaker. Claims *only* fall under this category if the quote to be verified directly figures in the claim, e.g. “*Boris Johnson told journalists ‘my favourite colour is red, because I love tomatoes’*”.
- **Position Statement.** The primary task is to identify whether a public figure has taken a certain position, e.g. supporting a particular policy or idea. For example, “*Edward Heath opposed privatisation*”. This also includes statements that opinions have changed, e.g. “*Edward Heath opposed privatisation before the election, but changed his mind after coming into office*”. Factual claims about the actions of people (e.g. “*Edward Heath nationalised Rolls-Royce*”) are not position statements (they are event or property claims); claims about the attitudes of people (e.g. “*Edward Heath supported the nationalisation of Rolls-Royce*”) are.
- **Event/Property Claim.** The primary task is to determine the veracity of a narrative about a particular event or series of events, or to identify whether a certain non-numerical property is true, e.g. a person attending a particular university. Some properties represent causal relationships, e.g. “*The prime minister never flies, because he has a fear of airplanes*”. In those cases, the claim should be interpreted as both a property claim and a causal claim.

- **Media Publishing Claim.** The primary task is to identify the original source for a (potentially doctored) image, video, or audio file. This covers both doctored media, and media that has been taken out of context (e.g. a politician is claimed to have shared a certain photo, and the task is to determine if they actually did). This also includes HTML-doctoring of social media posts. We will discard all claims in this category.
- **Media Analysis Claim.** The primary task is to perform complex reasoning about pieces of media, distinct from doctoring. This could for example be checking whether a geographical location is really where a video was taken, or determining whether a specific person is actually the speaker in an audio clip. The claim itself *must directly involve* media analysis; e.g. “the speaker of these two clips is the same”. Claims where the original source is video, but which can be understood and verified without viewing the original source, do not fall under this category. An original video or audio file can feature as metadata in fact-checking articles, but claims are only *complex media claims* if analysis of the video or audio beyond just extracting a quote is necessary for verification.

Several claim types – speculative claims, opinion claims, media publishing claims, and media analysis claims – will not be included in later phases.

3.5.3 Fact-checking Strategy

After identifying the claim type, we ask annotators to classify the approach taken by the fact checker according to the article. This is independent of the claim type, as a fact-checker might take any number of approaches to a given claim. Again, one *or several* options should be chosen from the following list:

- **Written Evidence.** The fact-checking process involved finding contradicting written evidence, e.g. a news article directly refuting the claim.
- **Numerical Comparison.** The fact-checking process involved numerical comparisons, such as verifying that one number is greater than another.
- **Consultation.** The fact checkers directly reached out to relevant experts or people involved with the story, reporting new information from such sources as part of the fact-checking article.
- **Satirical Source Identification.** The fact-checking process involved identifying the source of the claim as satire, e.g. The Onion.
- **Media Source Discovery.** The fact-checking process involved finding the original source of a (potentially doctored) image, video, or soundbite.
- **Image analysis.** The fact-checking process involved image analysis, such as comparing two images.
- **Video Analysis.** The fact-checking process involved analysing video, such as identifying the people in a video clip.
- **Audio Analysis** The fact-checking process involved analysing audio, such as determining which song was played in the background of an audio recording.
- **Geolocation.** The fact-checking process involved determining the geographical location of an image or a video clip, through the comparison of landmarks to pictures from Google Streetview or similar.
- **Fact-checker Reference.** The fact-checking process involved a reference to a previous fact-check of the same claim, either by the same or a different organisation. Reasoning or evidence from the referenced article was necessary to verify the claim.

Claims *only* labelled as solved through Fact-checker Reference will not be included in later phases.

4 Phase 2: Question Generation and Answering

The next round of annotation aims to produce pairs of questions and answers providing evidence to verify the claim. The primary sources of evidence are the URLs linked in the fact-checking article. We also provide access to a custom search bar to retrieve evidence.

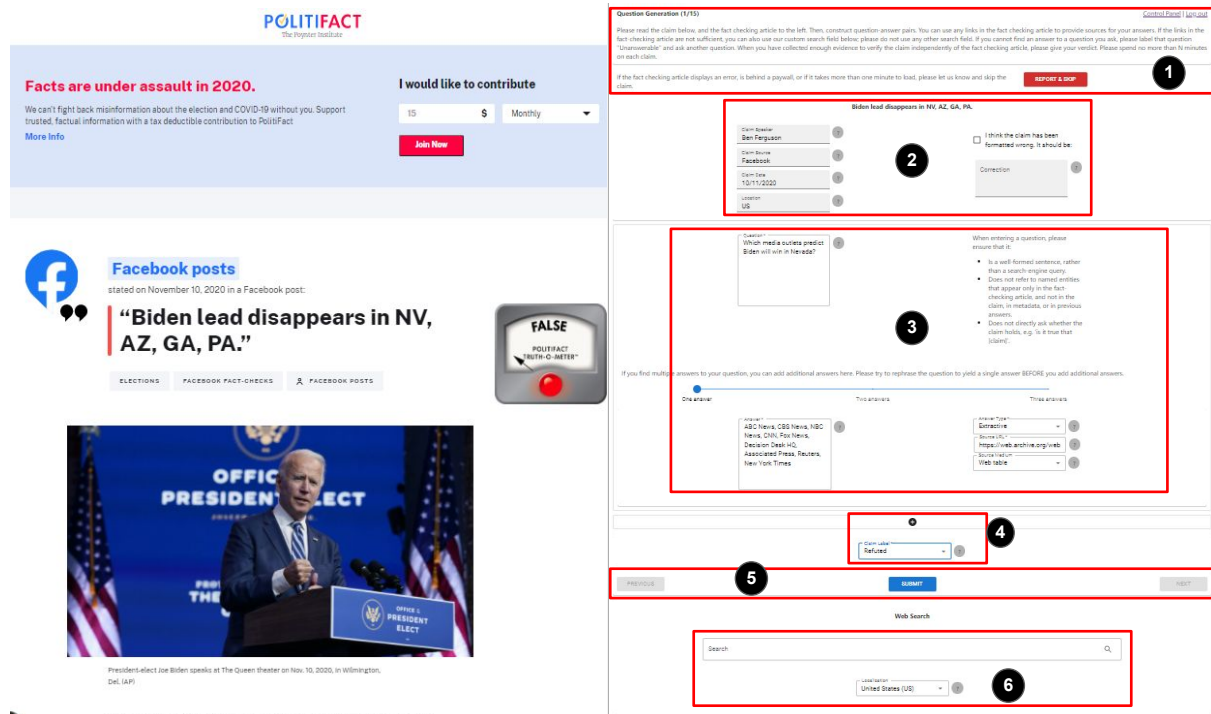


Figure 7: Interface of question generation. ① Guideline of annotation for this phase. Please read it before annotating. Notice that if the article displays a 404 page or another error, or if it takes more than one minute to load, please click the **REPORT & SKIP** button. ② The claim and the associated metadata. ③ Fields for the first question and its answers. Annotators can add up to 3 answers for each question if necessary. The text fields of metadata of question answer pairs are also provided. ④ Annotators can use the plus button to add as many questions as they want. Please select the label of this claim after finishing the question and answer generation. ⑤ Buttons for submitting the current claim, going to the previous claim, and next claim. ⑥ The custom search engine.

4.1 Overview

Here, we give a quick overview of the question generation task; in-detail discussion can be found in subsequent sections. Further documentation can also be found on-the-fly using the tooltips in the annotation interface.

1. The annotator should first read the claim and metadata provided by the previous annotator, and the associated fact-checking article (including the verdict). We note that because phase-one annotators sometimes split decompose claims into parts, in some cases not all sections of the fact-checking article will be relevant.
2. The task is then to generation questions and answers about the claim such that a verdict can be given without knowledge of the fact-checking article. The sources and strategies used in the fact-checking article can serve as inspiration for questions and evidence for answers, but the fact-checking article should not be *directly* referenced as a source.

3. We recommend constructing question-answer pairs iteratively, one at a time. That is, annotators should ask a question and attempt to answer it, and only then proceed to the next question.
4. Guidance on generating questions can be found in Section 4.2.
5. Answers should be sought from the metadata, any of the sources listed on the fact-checking article (e.g. any hyperlinks to other sites), and when that is not possible (e.g. due to the hyperlinks being dead) from the internet using the search bar we provide.
6. Questions about metadata can be used to draw attention to aspects of the claim, in order to reason about publication date or publication source (see Section 4.3).
7. **WARNING:** For persistence, we have stored all fact-checking articles on archive.org. Fact-checking articles may feature double-archived links using both archive.org and archive.is, e.g. "<https://web.archive.org/web/20201229212702/https://archive.md/28fMd>". Archive.org returns a 404 page for these. To view such a link, please just copy-paste the archive.is part (e.g. "<https://archive.md/28fMd>") into your browser.
8. Answers should be accompanied by a hyperlink to the source, and the type of the source – e.g. web text, a pdf – should be specified. We note that if the source type is set as metadata, the source link will automatically be set to the word *metadata*.
9. Answers can be either *extraction*, e.g. copy-pasted directly from the source, *abstractive*, e.g. written in free-form based on the source, or *boolean*, e.g. written as yes/no with an explanation taken either extractively or abstractively from the source. Where possible, we strongly prefer extractive answers.
10. If an answer cannot be found, we also allow annotators to mark the question as unanswerable. We ask annotators to use this instead of deleting unanswerable questions.
11. Guidance on generating answers can be found in Section 4.5.
12. If enough questions have been asked to support a verdict, or if at least ten minutes have passed without the annotator finding enough evidence, a verdict should be given from our for labels described in Section 3.4.
13. Annotators in phase two should base their verdict on the question-answer pairs they have generated, and *not* on the fact-checking article. Depending on what information has been retrieved, they may therefore disagree with the article.
14. Before proceeding to the next hit, the annotator will be shown a warning with the QA-pairs they have generated. They will also be shown their assigned label. They will be asked to confirm that the collected evidence is sufficient to assign the label they have chosen to the claim.
15. Sometimes, annotators may be in doubt as to whether an additional question should be added to further support the verdict. Generally speaking, we always prefer to have as many question-answer pairs as possible, so if in doubt annotators should veer on the side of adding that additional question.

Important! Annotators should not choose a label if the retrieved evidence does not support it; for example, if the label **conflicting evidence** is chosen, there should be evidence documenting the conflict. Labels in phase two can contradict the label of the fact-checker, if the annotator believes it is appropriate.

4.2 Question Generation

To ensure the quality of the generated questions, we ask the annotators to create their questions as follows:

- Questions should be well-formed, rather than search engine queries (e.g. “where is Cambridge?” rather than “Cambridge location”).
- Questions should be standalone and understandable without any previous questions.
- Questions should be based on the version of the claim shown in the interface (i.e. the version extracted by phase one annotators), and not on the version in the fact-checking article. If an annotator believes a phase one claim has been extracted wrongly, they can correct it using the appropriate box.
- The annotators should avoid any question that directly asks whether or not the claim holds, e.g. “*is it true that [claim]*”.
- The annotators should ask all questions necessary to gather the evidence needed for the verdict, including world knowledge that might seem obvious, but could depend for example on where one is from. For example, Europeans might have better knowledge of European geography/history than Americans, and vice-versa.
- As a guiding principle, at least 2 questions should be asked. This is not a hard limit, however, and the annotators can proceed with only one question asked if they do not feel more are needed.

The following are examples used to illustrate how questions should be asked. These are based on the real claim “*the US in 2017 has the largest percentage of immigrants, almost tied now with the historical high as a percentage of immigrants living in this country*”:

- Good: What was the population of the US in 2017?
- Good: How many immigrants live in the US in 2017?
- Bad: What was the population of the US? [No time specified to find a statistic]
- Bad: What was the population there in 2017? [What does *there* refer to?]
- Bad: Is it true that the US in 2017 has the largest percentage of immigrants, almost tied now with the historical high as a percentage of immigrants living in this country? [Directly paraphrases the claim]

4.3 Metadata

Questions about metadata can be used to draw attention to aspects of the claim, in order to reason about publication date or publication source. If, for example, the claim “*aliens made contact with earth March 3rd, 2021*” was published on September 1st, 2020, the publication date can be used to refute the claim. In such cases, we ask annotators to first generate a question/answer pair – “*when was this claim made?*” “*September 1st, 2020*” – which can then be used to refute the claim. Similarly, questions about publication source can be used to refute satirical claims – “*where was this claim published?*”, “*www.theonion.com*”, “*what is The Onion?*”, “*The Onion is an American digital media company and newspaper organization that publishes satirical articles on international, national, and local news.*”.

4.4 Common sense assumptions and world knowledge

As a part of the question generation process, annotators may have to make assumptions and/or use world knowledge to interpret the claim. For example, for the claim “*Shakira is Canadian*”, it may be necessary to choose what it means to be Canadian. This is expressed in how questions are formulated, e.g. “*does Shakira have Canadian citizenship?*” or “*where does Shakira live?*”. This may also involve politically charged judgments. For example, some First Nations people are classed as Canadian by the Canadian government, but do not use that label for themselves.

In such cases, we ask annotators to follow – as closely as possible – the judgments made by the fact-checking websites. If the annotators feel that these are incomplete or misleading, they can add additional questions.

For example, for the claim “*Edward Heath opposed privatisation*”, a fact checker might provide his party manifesto as evidence. A corresponding question could then be “*what did the 1970 Conservative Party manifesto say about privatisation?*” An annotator could encounter evidence for the nationalisation of Rolls Royce during Heath’s government, which the fact-checking article did not take into account. In that case, the annotator might want to add an additional question, such as “*did Heath’s government nationalise any companies?*”. The annotators should ask *both* questions.

Important! As opposed to Phase 1, annotators in Phase 2 *should* use their own judgment to assign labels (although they should not ignore evidence used by the fact-checker). As such, if they disagree with the fact-checker about the label, they can select a different label.

4.5 Answer Generation

To find answers to questions, the annotators can rely on metadata, or on any sources linked from the factchecking site. Where these fail to produce appropriate information – either because they are not relevant to an asked question or because they refer to sources which have been taken down – we provide search functionalities as an alternative. Note that the annotators are not allowed to use the fact-checking article itself as a source, only the pages *hyper-linked* in the fact-checking article (and only when they are not from fact-checking websites).

Once an answer has been found, annotators can choose between the following four options to enter it:

- **Extractive:** The answer can be copied directly from the source. We ask the annotators to use their browser’s copy-paste mechanism to enter it.
- **Abstractive:** A freeform answer can be constructed based on the source, but not directly copy-pasted.
- **Boolean:** This is a special case of abstractive answers, where a yes/no is sufficient to answer the question. A second box must be used to give an explanation for the verdict grounded in the source (e.g. “yes, because...”).
- **Unanswerable:** No source can be found to answer the question.

For extractive, abstractive, and boolean answers, the annotators are also asked to copy-paste a link to the source URL they used to answer the question. Extractive answers are preferred to abstractive and boolean answers.

In some cases, annotators might find different answers from different sources. Our annotation tools allows adding additional answers, up to three. While we provide this functionality, we ask that annotators try to rephrase the question to yield a single answer before adding additional answers.

We note that if the annotators can only find a *partial* answer to a question, they can still use that. In such cases, please give the partial answer rather than marking the question as unanswerable.

Our search engine marks pages originating from known sources of misinformation and/or satire. We do not prevent annotators from using such sources, but we ask that annotators avoid them if at all possible.

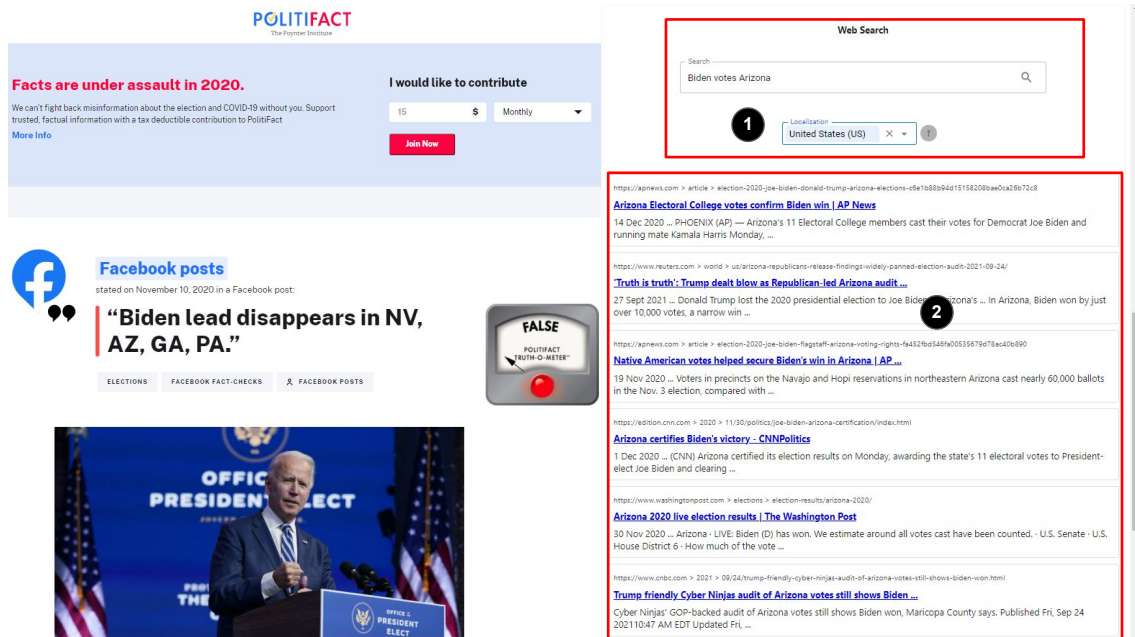


Figure 8: Interface of the search bar. ① Search bar and the location option. Annotators can change the localization of the search engine by selecting the country code here. ② Search results returned by the search engine.

In the event that an annotator wishes to use information from such a source, we strongly prefer that the finds similar, corroborating information from an additional source in order to further substantiate the evidence.

While answering a question, we furthermore ask annotators to adhere to the following:

Important!

- DO NOT use any other browser window/search bar to find an answer. You MUST use the provided search bar only.
- DO NOT give a verdict for the claim until you have finished questions and answers.
- DO NOT use the fact-checking article itself, or any other version of it you find on the internet, as evidence to support an answer.
- DO NOT submit answers using other articles from fact-checking websites, such as politifact.com or factcheck.org, as evidence.

4.6 Reasoning Chains of Claims

Annotators can build up reasoning chains across multiple questions, meaning that answers of one question can be used in the next question. For example, for the claim “*the fastest train in Japan drives at a top speed of 400 km/h*”, the first question is “What is the fastest Japanese train?”. The answer is “The fastest Japanese train is Shinkansen ALFA-X”. Based on the answer, we can further ask the second question to get more details, “What is the maximum operating speed of the Shinkansen ALFA-X”. Note that while the *generation* of the second question assumes knowledge of the answer to the first, it is *understandable* without it.

4.7 Confirmation

After submitting the question/answer pairs for a claim, annotators will be presented with a confirmation screen (see Figure 9). Annotators will be shown the question/answer pairs they have entered, along with

the verdict, and asked to confirm a second time that the verdict is supported by the evidence.

The screenshot shows a confirmation screen for a claim. At the top, the claim is: "Thorough hand-washing with an ordinary soap is effective in killing coronavirus (COVID-19)." Below this, there is a dropdown menu for "Claim Label" with "Supported" selected. To the right of the dropdown is a question mark icon. Below the dropdown are two buttons: "CANCEL" (purple) and "CONFIRM" (blue). To the right of the claim and label is a "Confirmation (3/20)" section with links for "Control Panel" and "Log out". Below this is a text box asking the user to confirm their verdict using only the question-answer pairs shown below. The interface is divided into two columns. The left column contains two question boxes: "How does soap kill coronavirus?" and "How does handwashing fight coronavirus?". The right column contains two answer boxes: "The outer layer of the virus is made up of lipids, aka fat. Soap dissolves that barrier, killing the virus." and "Coronavirus can spread through tiny droplets if the infected person coughs and sneezes. Touching any surface with droplets on, and then touching your eyes, nose, or mouth, can transmit the disease." Each answer box has a "View source" link below it.

Thorough hand-washing with an ordinary soap is effective in killing coronavirus (COVID-19).

Claim Label *
Supported

CANCEL CONFIRM

Confirmation (3/20) [Control Panel](#) | [Log out](#)

Please confirm that you can infer your chosen verdict using ONLY your question-answer pairs (shown below).

Question
How does soap kill coronavirus?

Answer
The outer layer of the virus is made up of lipids, aka fat. Soap dissolves that barrier, killing the virus.

[View source](#)

Question
How does handwashing fight coronavirus?

Answer
Coronavirus can spread through tiny droplets if the infected person coughs and sneezes. Touching any surface with droplets on, and then touching your eyes, nose, or mouth, can transmit the disease.

[View source](#)

Figure 9: Before moving on to the next claim, phase two annotators will be shown a confirmation screen to make sure that their chosen verdict is correct.

5 Phase 3: Quality Control

Once we have collected evidence in the form of generated questions and retrieved answers, we want to provide a measure of quality. Given a claim with associated evidence, we ask a third round of annotators to give a verdict for the claim. Crucially, the annotators at this round do not have access to the original fact-checking article, or to the claim label.

5.1 Overview

Here, we give a quick overview of the quality control task; in-detail discussion can be found in the following sections. Further documentation can also be found on-the-fly using the tooltips in the annotation interface.

1. Annotators should first read the claim, the metadata, and the question-answer pairs. This is the only information which should be used during this phase
2. It is important that annotators in the quality control phase do not use web search to find additional information, or rely on background knowledge which an average English speaker might not have. Commonsense facts that are known to (almost) everyone can be used – see Section 5.2.
3. If the claim, or any of the question-answer pairs lack context, they can be flagged. This helps us diagnose what is wrong with a set of question-answer pairs in the case annotators disagree over the label.
4. After reviewing the claim and the QA pairs, annotators should assign a label to the claim (see the four labels introduced in Section 3.4).
5. Finally, annotators should write a short statement justifying the verdict. If any commonsense information (e.g. background knowledge which an average English speaker *is* likely to have) is

used to give the verdict, but that information is not mentioned in any question-answer pair, it should be mentioned in the justification. For advice regarding justification production, see Section 5.3.

The screenshot displays a web interface for quality control. At the top, a header bar contains the title "Biden lead disappears in NV, AZ, GA, PA." and a "Control Panel | Log out" link. Below the header, there are several sections:

- Claim Speaker:** A dropdown menu showing "Facebook".
- Claim Source:** A dropdown menu showing "Facebook".
- Claim Date:** A dropdown menu showing "2020-11-10".
- Location:** A dropdown menu showing "US".
- Justification:** A text area containing the text: "Many media organizations believe Biden will win in NV, AZ, GA, and PA. As such, his lead has not disappeared."
- Claim Label:** A dropdown menu showing "Refuted".
- Verdict Validation (1/15):** A section with instructions: "Please read the claim and the question-answer pairs. Then, give your verdict on the claim. Do not look at any external information: make your verdict based ONLY on the question-answer pairs. If there are any problems with a question-answer pair, please use the form to report it. Do not use the information in any question-answer pair you report to make your verdict."

The main content area is divided into two columns. The left column contains a question: "Which media outlets predict Biden will win in Georgia?" and a section for "Potential problems with this question. Check any you think apply:" with checkboxes for "The question is not understandable/readable" and "The question is unrelated to the claim". The right column contains an answer: "No answer could be found." and a section for "Potential problems with this answer. Check any you think apply:" with checkboxes for "The answer is not understandable/readable", "Answer seems wrong, but is supported by the source", "Answer seems wrong, and is not supported by the source", and "I believe the source may be biased".

Below the main content area, there are two more question-answer pairs. The first question is "Which media outlets predict Biden will win in Arizona?" and the answer is "Fox News". The second question is "Which media outlets predict Biden will win in Georgia?" and the answer is "Associated Press".

At the bottom of the interface, there are three buttons: "PREVIOUS", "SUBMIT", and "NEXT".

Figure 10: Interface of quality control. ① Text field for entering the justification. ② Label of the claim and the checkbox of unreadable. Notice that once the unreadable option is selected, annotators do not need to select the label for the claim. ③ The question corresponds to the current claim. Here we have two question-answer pairs. If the annotator think the there exist potential problems with this question, check any options applied. ④ The answers corresponds to the question on the left. If the annotator think the there exist potential problems with the answer, check any options applied. ⑤ Buttons for submitting the current claim, going to the previous claim, and next claim.

5.2 Commonsense Knowledge

When giving a verdict, annotators sometimes need to rely on commonsense knowledge. Here, we consider only basic facts which an average English speaker is likely to know – e.g. “*Earth is a planet*” or “*raindrops consist of water*”. No other information beyond the question-answer pairs can be used in this phase.

We ask annotators to be relatively strict with what they consider commonsense, but use their own judgment. For example, we would consider “*Canada is a country*” commonsense, but not “*Canada is the third-largest country in terms of land mass*”. If an annotator is in doubt as to whether something is considered commonsense, they should not consider it commonsense.

5.3 Justification Production

In addition to the verdict, we as mentioned also ask annotators in Phase Three to write a short statement justifying their verdict. This justification should explain the reasoning process used to reach the verdict, along with any commonsense knowledge. If calculations or comparisons were used, e.g. “*6.3% is*

greater than 6.1%” or “ $10-4=6$ ”, they should be explicitly stated in the justification. Similarly, any rounding logic – e.g. “*4.3 million is approximately 4 million*” – should be explicitly stated here.

Other than commonsense knowledge, there should not be any new information presented in this statement. The justification should only describe how the annotators used the information present in the claim, the metadata, and the QA-pairs to reach their verdict. If a verdict cannot be reached, e.g. if the *not enough information*-label is chosen, annotators should instead describe what information is missing – e.g. “*I cannot determine if Canada is the third-largest country, because the questions do not specify how large any countries are.*”

6 Dispute Resolution

For some claims, there may be a disagreement between the labels produced by annotators in the question generation and quality control phases. In those cases, the claim will go through a second round of question generation and quality control. While the instructions given in Sections 4.2 and 5 still apply, we give a few extra recommendations specific to dispute resolution here.

6.1 Vague Claims

Some claims may pass to the dispute resolution phase because they are too vague for annotators in phases two and three to agree on the meaning. In order to catch these cases, the final step of dispute resolution – that is, the extra quality control step at the end – includes an additional label, *Claim Too Vague*. This should be select when and only when an annotator can understand the claim (e.g. it is readable), but there is too much doubt over how it is supposed to be interpreted. For example, the claim “*Ohio is the best state*” is too vague as it is not clear what “best” refers to.

6.2 Adding and Modifying Questions

The aim of dispute resolution is to resolve the conflict so that a potential new reader would come to a conclusive verdict. As such, the annotator should not necessarily agree with either the Phase Two or the Phase Three-verdict; they should attempt to make the fact-checking unambiguous. There may be cases where new questions must be added, and cases where existing questions should be changed but no new questions are necessary. There may also be cases where no change to the evidence is necessary at all, but where either the Phase Two or Phase Three-annotator has simply entered a wrong verdict. For this final category adding additional evidence to provide clarity can still be helpful, but it is not necessary; annotators should use their own judgment here.

6.3 NEI-verdicts

A common case for dispute resolution is the situation where the Phase Two annotator has selected *Supported*, *Refuted*, or *Conflicting Evidence/Cherrypicking* as the verdict, but the Phase Three annotator has selected *Not Enough Evidence*. This can happen for example if Phase Two annotators forget to gather some of the evidence they use to reach the verdict, rely on aspects only stated in the fact-checking article without making it explicit through a question-answer pair, or overestimate the strength of the evidence they have gathered. In these cases, the aim of dispute resolution is to gather additional evidence and resolve the conflict that way; i.e. it is not sufficient to give a *Not Enough Information*-verdict without attempting to add evidence (although the same time limit as in P2 applies).